



the
POWER
of
JAVA™



JavaOne
Sun and Microsoft are registered trademarks.

Simpler, Faster, Better: Concurrency Utilities in JDK™ Software Version 5.0

Brian Goetz

Principal Consultant, Quiotix Corp

David Holmes

Staff Engineer, Sun Microsystems Inc.

TS-4915

Goal

Learn how to use the new concurrency utilities (the **java.util.concurrent** package) to replace error-prone or inefficient code and to better structure applications

Agenda

Overview of `java.util.concurrent`

Concurrent Collections

Threads Pools and Task Scheduling

Locks, Conditions and Synchronizers

Atomic Variables

Rationale for `java.util.concurrent`

Developing concurrent classes was just too hard

- The built-in concurrency primitives – `wait()`, `notify()`, and `synchronized` – are, well, primitive
 - Hard to use correctly
 - Easy to use incorrectly
 - Specified at too low a level for most applications
 - Can lead to poor performance if used incorrectly
- Too much wheel-reinventing!

Goals for `java.util.concurrent`

Simplify development of concurrent applications

- Provide a set of basic concurrency building blocks
- Something for everyone:
 - Make some problems trivial to solve by everyone
 - Develop thread-safe classes, such as servlets, built on concurrent building blocks like `ConcurrentHashMap`
 - Make some problems easier to solve by concurrent programmers
 - Develop concurrent applications using thread pools, barriers, latches, and blocking queues
 - Make some problems possible to solve by concurrency experts
 - Develop custom locking classes, lock-free algorithms

Agenda

Overview of `java.util.concurrent`

Concurrent Collections

Threads Pools and Task Scheduling

Locks, Conditions and Synchronizers

Atomic Variables

Concurrent Collections

Concurrent vs. Synchronized

- Pre Java 5: *thread-safe* but not *concurrent* classes
- Thread-safe synchronized collections:
 - `Hashtable`, `Vector`, `Collections.synchronizedMap`
 - Monitor is source of contention under concurrent access
 - Often require locking during iteration
- Concurrent collections:
 - Allow multiple operations to overlap each other
 - Big performance advantage
 - At the cost of some slight differences in semantics
 - Might not support atomic operations

Concurrent Collections

- **ConcurrentHashMap**
 - Concurrent (scalable) replacement for **Hashtable** or **Collections.synchronizedMap**
 - Allows reads to overlap each other
 - Allows reads to overlap writes
 - Allows up to 16 writes to overlap
 - Iterators don't throw **ConcurrentModificationException**
- **CopyOnWriteArrayList**
 - Optimized for case where iteration is much more frequent than insertion or removal
 - Ideal for event listeners

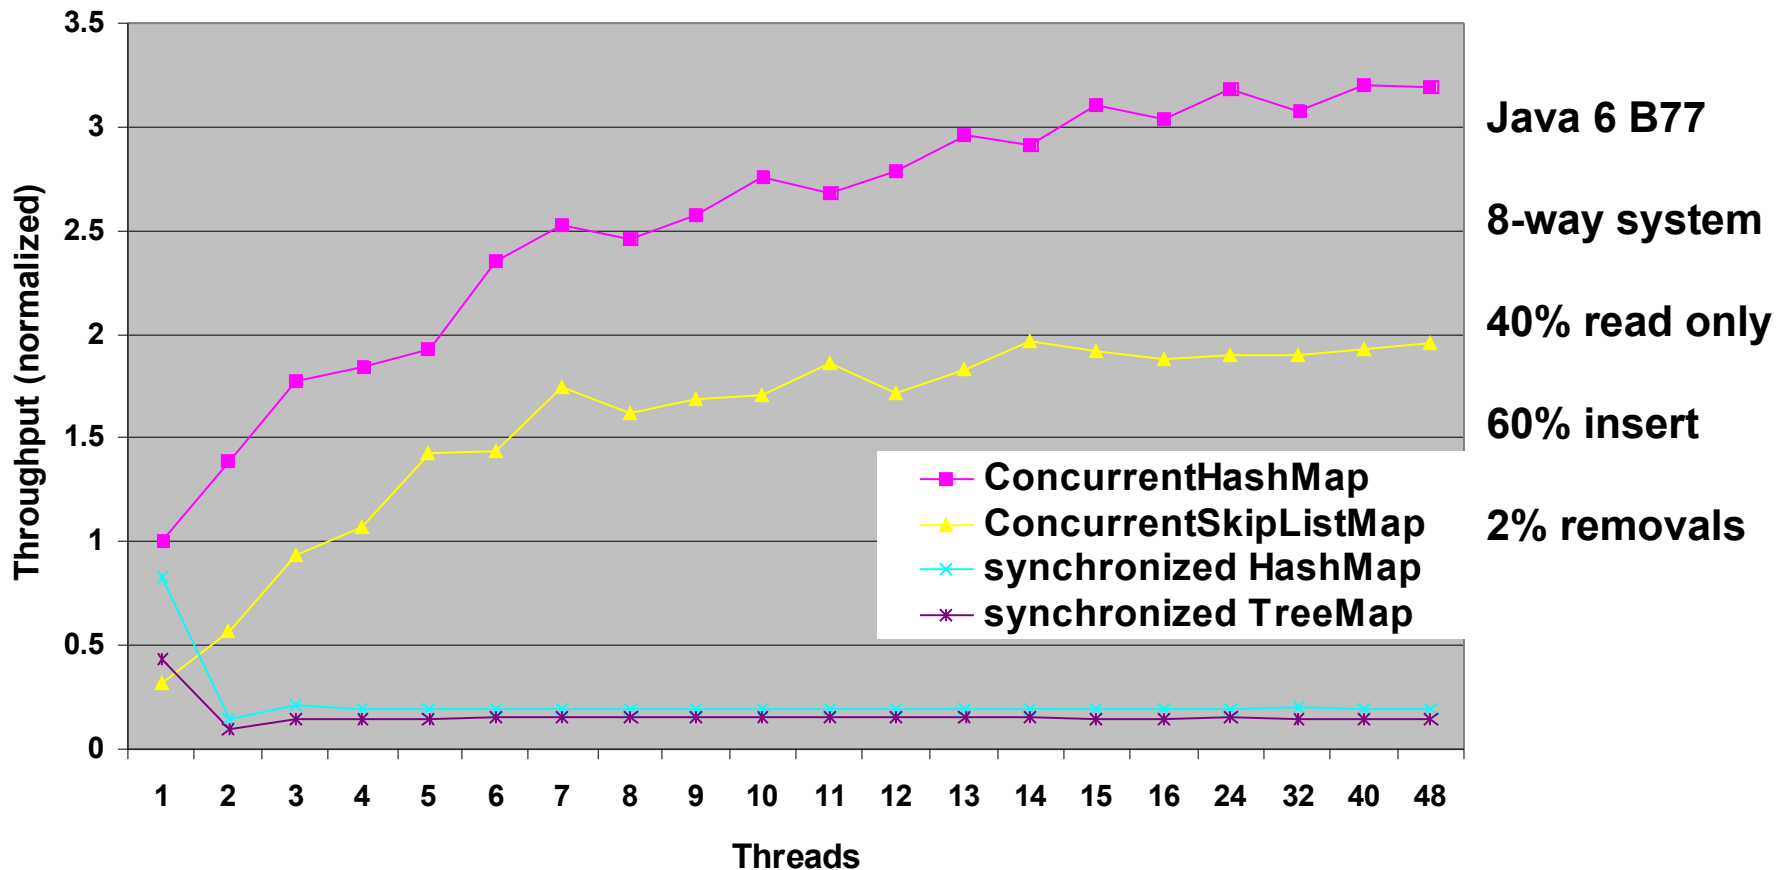
Concurrent Collections

Iteration Semantics

- Synchronized collection iteration broken by concurrent changes in another thread
 - Throws **ConcurrentModificationException**
 - Locking a collection during iteration hurts scalability
- Concurrent collections can be modified concurrently during iteration
 - Without locking the whole collection
 - Without **ConcurrentModificationException**
 - But changes may not be seen

Concurrent Collection Performance

Throughput in Thread-safe Maps



Queues

New Interface added to `java.util`

```
interface Queue<E> extends Collection<E> {
    boolean offer(E x);
    E poll();
    E remove() throws NoSuchElementException;
    E peek();
    E element() throws NoSuchElementException;
}
```

- Retrofit (non-thread-safe)—implemented by `LinkedList`
- Add (non-thread-safe) `PriorityQueue`
- Fast thread-safe non-blocking `ConcurrentLinkedQueue`
- Better performance than `LinkedList` is possible as random-access requirement has been removed

Blocking Queues

BlockingQueue Interface


- Extends **Queue** to provides blocking operations
 - Retrieval: **take** – wait for queue to become nonempty
 - Insertion: **put** – wait for capacity to become available
- Several implementations:
 - **LinkedBlockingQueue**
 - Ordered FIFO, may be bounded, two-lock algorithm
 - **PriorityBlockingQueue**
 - Unordered but retrieves least element, unbounded, lock-based
 - **ArrayBlockingQueue**
 - Ordered FIFO, bounded, lock-based
 - **SynchronousQueue**
 - Rendezvous channel, lock-based in Java 5, lock-free in Java 6

BlockingQueue Example

```
class LogWriter {
    final BlockingQueue msgQ =
        new LinkedBlockingQueue();

    public void writeMessage(String msg) throws IE {
        msgQ.put(msg);
    }

    // run in background thread
    public void logServer() {
        try {
            while (true) {
                System.out.println(msgQ.take());
            }
        }
        catch (InterruptedException ie) { ... }
    }
}
```



Producer-Consumer Pattern

- **LogWriter** example illustrates the *producer-consumer* pattern
 - Ubiquitous concurrency pattern, nearly always relies on some form of blocking queue
 - Decouples identification of work from doing the work
 - Simpler and more flexible
- **LogWriter** had many producers, one consumer
 - Thread pool has many producers, many consumers
- **LogWriter** moves IO from caller to log thread
 - Shorter code paths, fewer context switches, no contention for IO locks -> more efficient

Agenda

Overview of `java.util.concurrent`

Concurrent Collections

Threads Pools and Task Scheduling

Locks, Conditions and Synchronizers

Atomic Variables

Executors

Framework for asynchronous execution

- Standardize asynchronous invocation
 - Framework to execute `Runnable` and `Callable` tasks
- Separate submission from execution policy
 - Use `anExecutor.execute(aRunnable)`
 - Not `new Thread(aRunnable).start()`
- Cancellation and shutdown support
- Usually created via `Executors` factory class
 - Configures flexible `ThreadPoolExecutor`
 - Customize shutdown methods, before/after hooks, saturation policies, queuing

Executors

Decouple submission from execution policy

```
public interface Executor {  
    void execute(Runnable command);  
}
```

- Code which submits a task doesn't have to know in what thread the task will run
 - Could run in the calling thread, in a thread pool, in a single background thread (or even in another JVM!)
 - Executor implementation determines *execution policy*
 - Execution policy controls resource utilization, saturation policy, thread usage, logging, security, etc
 - Calling code need not know the execution policy

Executor and ExecutorService

ExecutorService adds lifecycle management

- **ExecutorService** supports both graceful and immediate shutdown

```
public interface ExecutorService extends Executor {  
    void shutdown();  
    List<Runnable> shutdownNow();  
    boolean isShutdown();  
    boolean isTerminated();  
    boolean awaitTermination(long time, TimeUnit unit)  
        throws InterruptedException;  
  
    // other convenience methods for submitting tasks  
}
```

- Many useful utility methods too

Creating Executors

Factory methods in the Executors class

```
public class Executors {  
    static ExecutorService  
        newSingleThreadedExecutor();  
  
    static ExecutorService  
        newFixedThreadPool(int poolSize);  
  
    static ExecutorService  
        newCachedThreadPool();  
  
    static ScheduledExecutorService  
        newScheduledThreadPool(int corePoolSize);  
  
    // additional versions specifying ThreadFactory  
    // additional utility methods  
}
```

Executors Example

Web Server – poor resource management

```
class UnstableWebServer {  
  
    public static void main(String[] args) {  
        ServerSocket socket = new ServerSocket(80);  
        while (true) {  
            final Socket connection = socket.accept();  
            Runnable r = new Runnable() {  
                public void run() {  
                    handleRequest(connection);  
                }  
            };  
            // Don't do this!  
            new Thread(r).start();  
        }  
    }  
}
```

Executors Example

Web Server – better resource management

```
class BetterWebServer {  
    Executor pool = Executors.newFixedThreadPool(7);  
  
    public static void main(String[] args) {  
        ServerSocket socket = new ServerSocket(80);  
        while (true) {  
            final Socket connection = socket.accept();  
            Runnable r = new Runnable() {  
                public void run() {  
                    handleRequest(connection);  
                }  
            };  
            pool.execute(r);  
        }  
    }  
}
```

Saturation Policies

- An Executor which execute tasks in a thread pool
 - Can guarantee you will not run out of threads
 - Can manage thread competition for CPU resources
- There is still a risk of running out of memory
 - Tasks could queue up without bound
- Solution: use a *bounded task queue*
 - Just so happens that JUC provides several of these...
- If queue fills up, the *saturation policy* is applied
 - Policies available: throw, discard oldest, discard newest, or run-in-calling-thread
 - The last has the benefit of throttling the load

Futures and Callables

Representing asynchronous tasks

- **Callable** is functional analog of **Runnable**

```
interface Callable<V> {  
    V call() throws Exception;  
}
```

- **Future** holds result of asynchronous call, normally a **Callable**

```
interface Future<V> {  
    V get() throws InterruptedException,  
        ExecutionException;  
    V get(long timeout, TimeUnit unit) throws ...;  
    boolean cancel(boolean mayInterrupt);  
    boolean isCancelled();  
    boolean isDone();  
}
```

Futures Example

Implementing a concurrent cache

```
public class Cache<K, V> {  
    final ConcurrentMap<K, FutureTask<V>> map =  
        new ConcurrentHashMap<K, FutureTask<V>>();  
  
    public V get(final K key) throws InterruptedException {  
        FutureTask<V> f = map.get(key);  
        if (f == null) {  
            Callable<V> c = new Callable<V>() {  
                public V call() {  
                    // return value associated with key  
                }  
            };  
            f = new FutureTask<V>(c);  
            FutureTask<V> old = map.putIfAbsent(key, f);  
            if (old == null)  
                f.run();  
            else  
                f = old;  
        }  
        try { return f.get(); }  
        catch (ExecutionException ex) { // rethrow ex.getCause() }  
    }  
}
```


ScheduledExecutorService

Deferred and recurring tasks

- **ScheduledExecutorService** can be used to:
 - Schedule a **Callable** or **Runnable** to run once with a fixed delay after submission
 - Schedule a **Runnable** to run periodically at a fixed rate
 - Schedule a **Runnable** to run periodically with a fixed delay between executions
- Submission returns a **ScheduledFutureTask** handle which can be used to cancel the task
- Like **java.util.Timer**, but supports pooling

Agenda

Overview of `java.util.concurrent`

Concurrent Collections

Threads Pools and Task Scheduling

Locks, Conditions and Synchronizers

Atomic Variables

Locks

- Use of monitor synchronization is just fine for most applications, but it has some shortcomings
 - Single wait-set per lock
 - No way to interrupt or time-out when waiting for a lock
 - Locking must be block-structured
 - Inconvenient to acquire a variable number of locks at once
 - Advanced techniques, such as hand-over-hand locking, are not possible
- Lock objects address these limitations
 - But harder to use: need **finally** block to ensure release
 - So if you don't need them, stick with **synchronized**

Framework for Flexible Locking

- ```
interface Lock {
 void lock();
 void lockInterruptibly() throws
 InterruptedException;
 boolean tryLock();
 boolean tryLock(long time, TimeUnit unit) throws
 InterruptedException;
 void unlock();
 Condition newCondition() throws
 UnsupportedOperationException;
}
```
- High-performance implementation: **ReentrantLock**
  - Basic semantics same as use of **synchronized**
  - Condition object semantics like **wait/notify**

# Simple Lock Example

- Used extensively within `java.util.concurrent`

```
final Lock lock = new ReentrantLock();
...
lock.lock();
try {
 // perform operations protected by lock
}
catch(Exception ex) {
 // restore invariants & rethrow
}
finally {
 lock.unlock();
}
```

- Must manually ensure lock is released

# Conditions

## Monitor-like operations for working with Locks

- **Condition** is an abstraction of **wait/notify**

```
interface Condition {
 void await() throws InterruptedException;
 boolean await(long time, TimeUnit unit)
 throws InterruptedException;
 long awaitNanos(long nanosTimeout)
 throws InterruptedException;
 boolean awaitUntil(Date deadline)
 throws InterruptedException;
 void awaitUninterruptibly();

 void signal();
 void signalAll();
}
```

- Timed **await** versions report reason for return

# Condition Example

```
class BoundedBuffer {
 final Lock lock = new ReentrantLock();
 final Condition notFull = lock.newCondition();
 final Condition notEmpty = lock.newCondition();
 ...
 void put(Object x) throws InterruptedException {
 lock.lock(); try {
 while (isFull()) notFull.await();
 doPut(x);
 notEmpty.signal();
 } finally { lock.unlock(); }
 }
 Object take() throws InterruptedException {
 lock.lock(); try {
 while (isEmpty()) notEmpty.await();
 notFull.signal();
 return doTake();
 } finally { lock.unlock(); }
 }
}
```

# Synchronizers

Utility classes for coordinating access and control

- **Semaphore** – Dijkstra counting semaphore, managing a specified number of permits
- **CountDownLatch** – allows one or more threads to wait for a set of threads to complete an action
- **CyclicBarrier** – allows a set of threads to wait until they all reach a specified barrier point
- **Exchanger** – allows two threads to rendezvous and exchange data
  - Such as exchanging an empty buffer for a full one



# Semaphore Example

Bound the submission of tasks to an executor

```
public class ExecutorProxy implements Executor {
 private final Semaphore tasks;
 private final Executor master;

 ExecutorProxy(Executor master, int limit) {
 this.master = master;
 tasks = new Semaphore(limit);
 }

 public void execute(Runnable r) {
 tasks.acquireUninterruptibly(); // for simplicity
 try {
 master.execute(r);
 }
 finally {
 tasks.release();
 }
 }
}
```

# Agenda

Overview of `java.util.concurrent`

Concurrent Collections

Threads Pools and Task Scheduling

Locks, Conditions and Synchronizers

**Atomic Variables**

# Atomic variables

Holder classes for scalars, references and fields

- Support atomic operations
  - Compare-and-set (CAS)
  - Get, set and arithmetic operations (where applicable)
    - Increment, decrement operations
- Abstraction of **volatile** variables
- Nine main classes:
  - { **int**, **long**, **reference** } X { **value**, **field**, **array** }
- E.g. **AtomicInteger** useful for counters, sequence numbers, statistics gathering

# AtomicInteger Example

Construction counter for monitoring/management

- Replace this:

```
class Service {
 static int services;
 public Service() {
 synchronized(Service.class) {
 services++;
 }
 } // ...
}
```
- With this:

```
class Service {
 static AtomicInteger services =
 new AtomicInteger();
 public Service() {
 services.getAndIncrement();
 }
 // ...
}
```

# Atomic Compare-and-Set (CAS)

- `boolean compareAndSet(int expected, int update)`
  - Atomically sets value to `update` if currently `expected`
  - Returns true on successful update
- Direct hardware support in all modern processors
  - CAS, cmpxchg, ll/sc
- High-performance on multi-processors
  - No locks, so no lock contention and no blocking
  - But can fail
    - So algorithms must implement retry loop
- Foundation of many concurrent algorithms

# Sneak Preview of Java 6

- Double-ended queues: `Deque`, `BlockingDeque`
  - Implementations: `ArrayDeque`,  
`LinkedBlockingDeque`, `ConcurrentLinkedDeque`
- Concurrent skiplists: `ConcurrentSkipList{Map|Set}`
- Enhanced navigation of sorted maps/sets
  - `Navigable{Map|Set}`
- Miscellaneous algorithmic enhancements
  - More use of lock-free algorithms in utilities
  - VM performance improvements for intrinsic locking
- M&M support for locks & conditions

# java.util.concurrent

- Executors
  - `Executor`
  - `ExecutorService`
  - `ScheduledExecutorService`
  - `Callable`
  - `Future`
  - `ScheduledFuture`
  - `Delayed`
  - `CompletionService`
  - `ThreadPoolExecutor`
  - `ScheduledThreadPoolExecutor`
  - `AbstractExecutorService`
  - `Executors`
  - `FutureTask`
  - `ExecutorCompletionService`
- Queues
  - `BlockingQueue`
  - `ConcurrentLinkedQueue`
  - `LinkedBlockingQueue`
  - `ArrayBlockingQueue`
  - `SynchronousQueue`
  - `PriorityBlockingQueue`
  - `DelayQueue`
- Concurrent Collections
  - `ConcurrentMap`
  - `ConcurrentHashMap`
  - `CopyOnWriteArray{List,Set}`
- Synchronizers
  - `CountDownLatch`
  - `Semaphore`
  - `Exchanger`
  - `CyclicBarrier`
- Locks: `java.util.concurrent.locks`
  - `Lock`
  - `Condition`
  - `ReadWriteLock`
  - `AbstractQueuedSynchronizer`
  - `LockSupport`
  - `ReentrantLock`
  - `ReentrantReadWriteLock`
- Atomics: `java.util.concurrent.atomic`
  - `Atomic{Type}`
  - `Atomic{Type}Array`
  - `Atomic{Type}FieldUpdater`
  - `Atomic{Markable,Stampable}Reference`

# Summary

- Whenever you are about to use
  - `Object.wait`, `notify`, `notifyAll`
  - `new Thread(aRunnable).start()` ;
  - `synchronized`
- Check first in `java.util.concurrent` if there is a class that ...
  - Does it already, or
  - Let's you do it a simpler, or better way, or
  - Provides a better starting point for your own solution
- Don't reinvent the wheel!



# For More Information

- JavaDoc for `java.util.concurrent` – in JDK 5.0 download or on Sun website
- Doug Lea's concurrency-interest mailing list
  - <http://gee.cs.oswego.edu/dl/concurrency-interest/index.html>
- Concurrent Programming in Java (Lea)
  - Addison-Wesley, 1999 ISBN 0-201-31009-0
- Java Concurrency in Practice (Goetz, et al)
  - Addison-Wesley, 2006, ISBN 0-321-34960-1
- JUC Backport to JDK 1.4
  - <http://www.mathcs.emory.edu/dcl/util/backport-util-concurrent/>

# Q&A

Brian Goetz

David Holmes



the  
**POWER**  
of  
**JAVA™**



**JavaOne**  
Sun and Microsoft are leading partners

# Simpler, Faster, Better: Concurrency Utilities in JDK™ Software Version 5.0

**Brian Goetz**

Principal Consultant, Quotix Corp

**David Holmes**

Staff Engineer, Sun Microsystems Inc.

TS-4915