

# Difference between Dirty Read, Non Repeatable Read and Phantom Read in Database

## Dirty Read:-

Dirty read occurs when one transaction is changing the record, and the other transaction can read this record before the first transaction has been committed or rolled back. This is known as a dirty read scenario because there is always the possibility that the first transaction may rollback the change, resulting in the second transaction having read an invalid data.

## Dirty Read Example:-

Transaction A begins.  
UPDATE EMPLOYEE SET SALARY = 10000 WHERE EMP\_ID= '123';

Transaction B begins.  
SELECT \* FROM EMPLOYEE;  
(Transaction B sees data which is updated by transaction A. But, those updates have not yet been committed.)

## Non-Repeatable Read:-

Non Repeatable Reads happen when in a same transaction same query yields to a different result. This occurs when one transaction repeatedly retrieves the data, while a difference transactions alters the underlying data. This causes the different or non-repeatable results to be read by the first transaction.

## Non-Repeatable Example:-

Transaction A begins.  
SELECT \* FROM EMPLOYEE WHERE EMP\_ID= '123';

Transaction B begins.  
UPDATE EMPLOYEE SET SALARY = 20000 WHERE EMP\_ID= '123';  
(Transaction B updates rows viewed by the transaction A before transaction A commits.) If Transaction A issues the same SELECT statement, the results will be different.

## Phantom Read:-

Phantom read occurs where in a transaction execute same query more than once, and the second transaction result set includes rows that were not visible in the first result set. This is caused by another transaction inserting new rows between the execution of the two queries. This is similar to a non-repeatable read, except that the number of rows is changed either by insertion or by deletion.

## Phantom Read Example:-

Transaction A begins.  
SELECT \* FROM EMPLOYEE WHERE SALARY > 10000 ;

Transaction B begins.  
INSERT INTO EMPLOYEE (EMP\_ID, FIRST\_NAME, DEPT\_ID, SALARY) VALUES ('111', 'Jamie', 10, 35000);  
Transaction B inserts a row that would satisfy the query in Transaction A if it were issued again.



## Related Posts:

1. [Difference between Optimistic Locking & Pessimistic Locking in Database](#)
2. [Isolation Levels in Database](#)
3. [Displaying Even and Odd number of Records in Database](#)
4. [How to eliminate the duplicate rows in Database](#)
5. [How to Retrieve Nth Maximum Salary in Database](#)
6. [Performance and Tuning Tips in Database](#)

March 9th, 2015 | Tags: [SQL](#) | Category: [Database, Java](#)

# DIRTY READS , PHANTOM READS

**DIRTY READS**: Reading uncommitted modifications are call Dirty Reads. Values in the data can be changed and rows can appear or disappear in the data set before the end of the transaction, thus getting you incorrect or wrong data.

This happens at **READ UNCOMMITTED** transaction isolation level, the lowest level. Here transactions running do not issue **SHARED locks** to prevent other transactions from modifying data read by the current transaction. This also do not prevent from reading rows that have been modified but not yet committed by other transactions.

To prevent Dirty Reads, READ COMMITTED or SNAPSHOT isolation level should be used.

**PHANTOM READS**: Data getting changed in current transaction by other transactions is called Phantom Reads. New rows can be added by other transactions, so you get different number of rows by firing same query in current transaction.

In **REPEATABLE READ** isolation levels Shared locks are acquired. This prevents data modification when other transaction is reading the rows and also prevents data read when other transaction are modifying the rows. But this does not stop INSERT operation which can add records to a table getting modified or read on another transaction. This leads to PHANTOM reads.

PHANTOM reads can be prevented by using **SERIALIZABLE** isolation level, the highest level. This level acquires **RANGE locks** thus preventing READ, Modification and INSERT operation on other transaction until the first transaction gets completed.

**DIRTY READS**: Reading uncommitted modifications are call Dirty Reads. Values in the data can be changed and rows can appear or disappear in the data set before the end of the transaction, thus getting you incorrect or wrong data.

This happens at **READ UNCOMMITTED** transaction isolation level, the lowest level. Here transactions running do not issue **SHARED locks** to prevent other transactions from modifying data read by the current transaction. This also do not prevent from reading rows that have been modified but not yet committed by other transactions.

To prevent Dirty Reads, READ COMMITTED or SNAPSHOT isolation level should be used.

**PHANTOM READS**: Data getting changed in current transaction by other transactions is called Phantom Reads. New rows can be added by other transactions, so you get different number of rows by firing same query in current transaction.

In **REPEATABLE READ** isolation levels Shared locks are acquired. This prevents data modification when other transaction is reading the rows and also prevents data read when other transaction are modifying the rows. But this does not stop INSERT operation which can add records to a table getting modified or read on another transaction. This leads to PHANTOM reads.

PHANTOM reads can be prevented by using **SERIALIZABLE** isolation level, the highest level. This level acquires **RANGE locks** thus preventing READ, Modification and INSERT operation on other transaction until the first transaction gets completed.

# Spring @Transactional - isolation, propagation

**PROPAGATION\_REQUIRED = 0;** If DataSourceTransactionObject T1 is already started for Method M1.If for another Method M2 Transaction object is required ,no new Transaction object is created .Same object T1 is used for M2

**PROPAGATION\_MANDATORY = 2;** method must run within a transaction. If no existing transaction is in progress, an exception will be thrown

**PROPAGATION\_REQUIRES\_NEW = 3;** If DataSourceTransactionObject T1 is already started for Method M1 and it is in progress(executing method M1) .If another method M2 start executing then T1 is suspended for the duration of method M2 with new DataSourceTransactionObject T2 for M2.M2 run within its own transaction context

**PROPAGATION\_NOT\_SUPPORTED = 4;** If DataSourceTransactionObject T1 is already started for Method M1.If another method M2 is run concurrently .Then M2 should not run within transaction context. T1 is suspended till M2 is finished.

**PROPAGATION\_NEVER = 5;** None of the methods run in transaction context.

**An isolation level:** It is about how much a transaction may be impacted by the activities of other concurrent transactions.It supports consistency leaving the data across many tables in a consistent state. It involves locking rows and/or tables in a database.

## The problem with multiple transaction

**Scenario 1.**If T1 transaction reads data from table A1 that was written by another concurrent transaction T2.If on the way T2 is rollback,the data obtained by T1 is invalid one.E.g a=2 is original data .If T1 read a=1 that was written by T2.If T2 rollback then a=1 will be rollback to a=2 in DB.But,Now ,T1 has a=1 but in DB table it is changed to a=2.

**Scenario2.**If T1 transaction reads data from table A1.If another concurrent transaction(T2) update data on table A1.Then the data that T1 has read is different from table A1.Because T2 has updated the data on table A1.E.g if T1 read a=1 and T2 updated a=2.Then a!=b.

**Scenario 3.**If T1 transaction reads data from table A1 with certain number of rows. If another concurrent transaction(T2) inserts more rows on table A1.The number of rows read by T1 is different from rows on table A1

Scenario 1 is called **Dirty reads**

Scenario 2 is called **Nonrepeatable reads**

Scenario 3 is called **Phantom reads** .

So,isolation level is the extend to which **Scenario 1 ,Scenario 2 ,Scenario 3** can be prevented. You can obtained complete isolation level by implementing locking.That is preventing concurrent reads and writes to the same data from occurring.But it affects performance .The level of isolation depends upon application to application how much isolation is required.

**ISOLATION\_READ\_UNCOMMITTED** :Allows to read changes that haven't yet been committed.It suffer from Scenario 1 ,Scenario 2 ,Scenario 3

**ISOLATION\_READ\_COMMITTED**:Allows reads from concurrent transactions that have been committed.It may suffer from Scenario 2 ,Scenario 3 . Because other transactions may be updating the data.

**ISOLATION\_REPEATABLE\_READ**:Multiple reads of the same field will yield the same results until it is changed by itself.It may suffer from Scenario 3.Because other transactions may be inserting the data

**ISOLATION\_SERIALIZABLE**: Scenario 1,Scenario 2,Scenario 3 never happens.It is complete isolation.It involves full locking.It affects performance because of locking.

You can test using

```
public class TransactionBehaviour {  
    // set is either using xml Or annotation  
    DataSourceTransactionManager manager=new DataSourceTransactionManager();  
    SimpleTransactionStatus status=new SimpleTransactionStatus();  
};
```

```

public void beginTransaction()
{
    DefaultTransactionDefinition Def = new DefaultTransactionDefinition();
    // overwrite default PROPAGATION_REQUIRED and ISOLATION_DEFAULT
    // set is either using xml Or annotation
    manager.setPropagationBehavior(XX);
    manager.setIsolationLevelName(XX);

    status = manager.getTransaction(Def);
}

public void commitTransaction()
{
    if(status.isCompleted()){
        manager.commit(status);
    }
}

public void rollbackTransaction()
{
    if(!status.isCompleted()){
        manager.rollback(status);
    }
}

Main method{
    beginTransaction()
    M1();
    If error(){
        rollbackTransaction()
    }
    commitTransaction();
}
}

```

You can debug and see the result with different values for isolation and propagation.

Good question, although not a trivial one to answer.

## Propagation

Defines how transactions relate to each other. Common options

- **Required**: Code will always run in a transaction. Create a new transaction or reuse one if available.
- **Requires\_new**: Code will always run in a new transaction. Suspend current transaction if one exist.

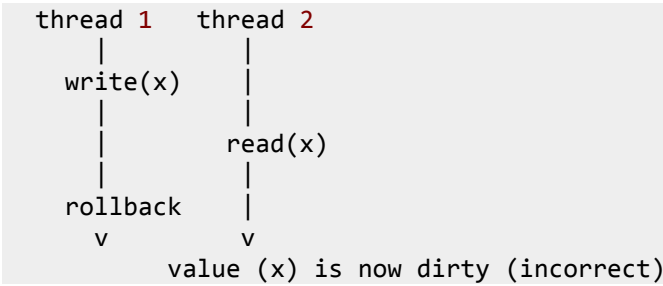
## Isolation

Defines the data contract between transactions.

- **Read Uncommitted**: Allows dirty reads
- **Read Committed**: Does not allow dirty reads
- **Repeatable Read**: If a row is read twice in the same transaciton, result will always be the same
- **Serializable**: Performs all transactions in a sequence

The different levels have different performance characteristics in a multi threaded application. I think if you understand the **dirty reads** concept you will be able to select a good option.

## Example when a dirty read can occur



So a sane default (if such can be claimed) could be `Read Committed`, which only lets you read values which have already been committed by other running transactions, in combination with an isolation level of `Required`. Then you can work from there if your application has other needs.

A practical example where a new transaction will always be created when entering the `provideService` routine and completed when leaving.

```
public class FooService {
    private Repository repo1;
    private Repository repo2;

    @Transactional(propagation=Propagation.REQUIRES_NEW)
    public void provideService() {
        repo1.retrieveFoo();
        repo2.retrieveFoo();
    }
}
```

Had we used `Required` instead the transaction **will remain open** if the transaction was already open when entering the routine. Note also that the result of a `rollback` could be different as several executions could take part in the same transaction.

We can easily verify the behaviour with a test and see how results differ with propagation levels

```
@RunWith(SpringJUnit4ClassRunner.class)
@ContextConfiguration(locations="classpath:/fooService.xml")
public class FooServiceTests {

    private @Autowired TransactionManager transactionManager;
    private @Autowired FooService fooService;

    @Test
    public void testProvideService() {
        TransactionStatus status = transactionManager.getTransaction(new
DefaultTransactionDefinition());
        fooService.provideService();
        transactionManager.rollback(status);
        // assert repository values are unchanged ...
    }
}
```

With a propagation level of

- `Requires new` we would expect `fooService.provideService()` was *NOT* rolled back since it created its own sub-transaction.
- `Required` we would expect everything was rolled back and backing store unchanged.

Enough explanation about each parameter is given by other answers; However you asked for a real world example, here is the one that clarifies the purpose of different **propagation** options: Suppose you're in charge of implementing a *signup service* in which a confirmation e-mail is sent to the user. You come up with two service objects, one for *enrolling* the user and one for *sending* e-mails, which the latter is called inside the first one. For example something like this:

```
/* Sign Up service */
@Service
@Transactional(Propagation=REQUIRED)
class SignUpService{
    ...
    void SignUp(User user){
        ...
        SendEmail(User);
    }
}

/* E-Mail Service */
@Service
@Transactional(Propagation=REQUIRES_NEW)
class EmailService{
    ...
    void sendMail(User user){
        try{
            ... // Trying to send the e-mail
        }catch( Exception)
        {}
    }
}
```

You may have noticed that the second service is of propagation type **REQUIRES\_NEW** and moreover chances are it throws an exception (SMTP server down, invalid e-mail or other reasons). You probably don't want the whole process to roll-back, like removing the user information from database or other things; therefore you call the second service in a separate transaction.

Back to our example, this time you are concerned about the database security, so you define your DAO classes this way:

```
/* User DAO */
@Transactional(Propagation=MANDATORY)
class UserDAO{
    // some CRUD methods
}
```

Meaning that whenever a DAO object, and hence a potential access to db, is created, we need to reassure that the call was made from inside one of our services, implying that a live transaction should exist; otherwise an exception occurs. Therefore the propagation is of type **MANDATORY**.

---

You almost never want to use `Read Uncommitted` since it's not really `ACID` compliant. `Read Committed` is a good default starting place. `Repeatable Read` is probably only needed in reporting, rollup or aggregation scenarios. Note that many DBs, postgres included don't actually support `Repeatable Read`, you have to use `Serializable` instead. `Serializable` is useful for things that you know have to happen completely independently of anything else; think of it like `synchronized` in Java. `Serializable` goes hand in hand with `REQUIRES_NEW` propagation.

I use `REQUIRES` for all functions that run `UPDATE` or `DELETE` queries as well as "service" level functions. For DAO level functions that only run `SELECTs`, I use `SUPPORTS` which will participate in a TX if one is already started (i.e. being called from a service function).

---

Transaction Isolation and Transaction Propagation although related but are clearly two very different concepts. In both cases defaults are customized at client boundary component either by using [Declarative transaction management](#) or [Programmatic transaction management](#). Details of each isolation levels and propagation attributes can be found in reference links below.

**Transaction Isolation**

For given two or more running transactions/connections to a database, how and when are changes made by queries in one transaction impact/visible to the queries in a different transaction. It also related to what kind of database record locking will be used to isolate changes in this transaction from other transactions and vice versa. This is typically implemented by database/resource that is participating in transaction.

## Transaction Propagation

In an enterprise application for any given request/processing there are many components that are involved to get the job done. Some of this components mark the boundaries (start/end) of a transaction that will be used in respective component and it's sub components. For this transactional boundary of components, Transaction Propagation specifies if respective component will or will not participate in transaction and what happens if calling component already has or does not have a transaction already created/started. This is same as Java EE Transaction Attributes. This is typically implemented by the client transaction/connection manager.

---

---

## Spring Transaction Attributes

### What are transaction attributes?

Spring transactions allow setting up the propagation behavior, isolation, timeout and read only settings of a transaction. Before we delve into the details, here are some points that need to be kept in mind

- Isolation level and timeout settings get applied only after the transaction starts.
- Not all transaction managers specify all values and may throw exception with some non default values

### Propagation

#### PROPAGATION\_REQUIRED

This attribute tells that the code needs to be run in a transactional context. If a transaction already exists then the code will use it otherwise a new transaction is created. This is the default and mostly widely used transaction setting.

#### PROPAGATION\_SUPPORTS

If a transaction exists then the code will use it, but the code does not require a new one. As an example, consider a ticket reservation system. A query to get total seats available can be executed non-transactionally. However, if used within a transaction context it will deduct tickets already selected and reduce them from the total count, and hence may give a better picture. This attribute should be used with care especially when PROPAGATION\_REQUIRED or PROPAGATION\_REQUIRES\_NEW is used within a PROPAGATION\_SUPPORTS context.

#### PROPAGATION\_MANDATORY

Participates in an existing transaction, however if no transaction context is present then it throws a TransactionRequiredException

#### PROPAGATION\_REQUIRES\_NEW

Creates a new transaction and if an existing transaction is present then it is suspended. In other words a new transaction is always started. When the new transaction is complete then the original transaction resumes. This transaction type is useful when a sub activity needs to be completed irrespective of the containing transaction. The best example of this is logging. Even if a transaction roll backs you still want to preserve the log statements. Transaction suspension may not work out of the box with all transaction managers, so make sure that the transaction manager supports transaction suspension



## PROPAGATION\_NOT\_SUPPORTED

This attribute says that transaction is not supported. In other words the activity needs to be performed non-transactionally. If an existing transaction is present then it is suspended till the activity finishes.

## PROPAGATION\_NEVER

This attribute says that the code cannot be invoked within a transaction. However, unlike PROPAGATION\_NOT\_SUPPORTED, if an existing transaction is present then an exception will be thrown

## PROPAGATION\_NESTED

The code is executed within a nested transaction if existing transaction is present, if no transaction is present then a new transaction is created. Nested transaction is supported out of the box on only certain transaction managers.

## Isolation

Isolation is a property of a transaction that determines what effect a transaction has on other concurrent transactions. To completely isolate the transaction the database may apply locks to rows or tables. Before we go through the transaction levels, let us look at some problems that occur when transaction 1 reads data that is being modified by transaction 2.

- *Dirty Reads*- Dirty reads occur when transaction 2 reads data that has been modified by transaction 1 but not committed. The problem occurs when transaction 1 rolls back the transaction, in which case the data read by transaction 2 will be invalid.
- *Non Repeatable Reads*- Nonrepeatable reads happen when a transaction fires the same query multiple times but receives different data each time for the same query. This may happen when another transaction has modified the rows while this query is in progress.
- *Phantom Reads* - Phantom reads occur when the collection of rows returned is different when a same query is executed multiple times in a transaction. Phantom reads occur when transaction 2 adds rows to a table between the multiple queries of transaction 1.

The following isolation levels are supported by spring

## ISOLATION\_DEFAULT

Use the isolation level of the underlying database.

## ISOLATION\_READ\_UNCOMMITTED

This is the lowest level of isolation and says that a transaction is allowed to read rows that have been added but not committed by another transaction. This level allows dirty reads, phantom reads and non repeatable reads.

## ISOLATION\_READ\_COMMITTED

This level allows multiple transactions on the same data but does not allow uncommitted transaction of one transaction to be read by another. This level, therefore, prevents dirty reads but allows phantom reads and nonrepeatable reads. This is the default isolation setting for most database and is supported by most databases.

## ISOLATION\_REPEATABLE\_READ

This level ensures that the data set read during a transaction remains constant even if another transaction modifies and commits changes to the data. Therefore if transaction 1 reads 4 rows of data and transaction 2 modifies and commits the fourth row and then transaction 1 reads the four rows again then it does not see the modifications made by transaction 2. (It does not see the changes made in the fourth row by the second transaction). This level prevents dirty reads and non repeatable reads but allows phantom reads.

## ISOLATION\_SERIALIZABLE

This is the highest isolation level. It prevents dirty reads, non repeatable reads and phantom reads. This level prevents the situation when transaction 1 performs a query with a certain where clause and retrieves say four rows, transaction 2 inserts a row that forms part of the same where clause and then transaction 1 reruns the query with the same where clause but still sees only four rows (does not see the row added by the second transaction)

## **Read Only**

The read only attribute specifies that the transaction is only going to read data from a database. The advantage is that the database may apply certain optimization to the transaction when it is declared to be read only. Since read only attribute comes in action as soon as the transaction starts, it may be applied to only those propagation settings that start a transaction. i.e. PROPAGATION\_REQUIRED, PROPAGATION\_REQUIRES\_NEW and PROPAGATION\_NESTED.

## **Timeout**

Timeout specifies the maximum time allowed for a transaction to run. This may be required since transactions that run for a very long time may unnecessarily hold locks for a long time. When a transaction reaches the timeout period, it is rolled back. Timeout needs to be specified only on propagation settings that start a new transaction

## **Rollback Rules**

It is also possible to specify that transactions roll back on certain exceptions and do not rollback on other exceptions by specifying the rollback rules.

# Isolation (database systems)

From Wikipedia, the free encyclopedia

In database systems, **isolation** determines how transaction integrity is visible to other users and systems. For example, when a user is creating a Purchase Order and has created the header, but not the PO lines, is the header available for other systems/users, carrying out concurrent operations (such as a report on Purchase Orders), to see?

A lower isolation level increases the ability of many users to access data at the same time, but increases the number of concurrency effects (such as dirty reads or lost updates) users might encounter. Conversely, a higher isolation level reduces the types of concurrency effects that users may encounter, but requires more system resources and increases the chances that one transaction will block another.<sup>[1]</sup>

It is typically defined at database level as a property that defines how/when the changes made by one operation become visible to other, but on older systems may be implemented systemically, for example through the use of temporary tables. In two-tier systems, a TP (Transaction Processing) manager is required to maintain isolation. In n-tier systems (such as multiple websites attempting to book the last seat on a flight) a combination of stored procedures and transaction management is required to commit the booking and confirm to the customer.<sup>[2]</sup>

Isolation is one of the ACID (Atomicity, Consistency, Isolation, Durability) properties.

## Contents

- 1 Concurrency control
- 2 Isolation levels
  - 2.1 Serializable
  - 2.2 Repeatable reads
  - 2.3 Read committed
  - 2.4 Read uncommitted
- 3 Default isolation level
- 4 Read phenomena
  - 4.1 Dirty reads
  - 4.2 Non-repeatable reads
  - 4.3 Phantom reads
- 5 Isolation Levels, Read Phenomena and Locks
  - 5.1 Isolation Levels vs Read Phenomena
  - 5.2 Isolation Levels vs Lock Duration
- 6 See also
- 7 References
- 8 External links

## Concurrency control

Concurrency control comprises the underlying mechanisms in a DBMS which handles isolation and guarantees related correctness. It is heavily utilized by the database and storage engines (see above) both to guarantee the correct execution of concurrent transactions, and (different mechanisms) the correctness of other DBMS processes. The transaction-related mechanisms typically constrain the database data access operations' timing (transaction schedules) to certain orders characterized as the serializability and recoverability schedule properties. Constraining database access operation execution typically means reduced performance (rates of execution), and thus concurrency control mechanisms are typically designed to provide the best performance possible under the constraints. Often, when possible without harming correctness, the serializability property is compromised for better performance. However, recoverability cannot be compromised, since such typically results in a quick database integrity violation.

Two-phase locking is the most common transaction concurrency control method in DBMSs, used to provide both serializability and recoverability for correctness. In order to access a database object a transaction first needs to acquire a lock for this object. Depending on the access operation type (e.g., reading or writing an object) and on the lock type, acquiring the lock may be blocked and postponed, if another transaction is holding a lock for that object.

## Isolation levels

Of the four ACID properties in a DBMS (Database Management System), the isolation property is the one most often relaxed. When attempting to maintain the highest level of isolation, a DBMS usually acquires locks on data or implements multiversion concurrency control, which may result in a loss of concurrency. This requires adding logic for the application to function correctly.

Most DBMSs offer a number of *transaction isolation levels*, which control the degree of locking that occurs when selecting data. For many database applications, the majority of database transactions can be constructed to avoid requiring high isolation levels (e.g. `SERIALIZABLE` level), thus reducing the locking overhead for the system. The programmer must carefully analyze database access code to ensure that any relaxation of isolation does not cause software bugs that are difficult to find. Conversely, if higher isolation levels are used, the possibility of deadlock is increased, which also requires careful analysis and programming techniques to avoid.

The isolation levels defined by the ANSI/ISO SQL standard are listed as follows.

### Serializable

This is the *highest* isolation level.

With a lock-based concurrency control DBMS implementation, serializability requires read and write locks (acquired on selected data) to be released at the end of the transaction. Also *range-locks* must be acquired when a `SELECT` query uses a ranged *WHERE* clause, especially to avoid the *phantom reads* phenomenon (see below).

When using non-lock based concurrency control, no locks are acquired; however, if the system detects a *write collision* among several concurrent transactions, only one of them is allowed to commit. See *snapshot isolation* for more details on this topic.

### Repeatable reads

In this isolation level, a lock-based concurrency control DBMS implementation keeps read and write locks (acquired on selected data) until the end of the transaction. However, *range-locks* are not managed, so *phantom reads* can occur.

## Read committed

In this isolation level, a lock-based concurrency control DBMS implementation keeps write locks (acquired on selected data) until the end of the transaction, but read locks are released as soon as the SELECT operation is performed (so the *non-repeatable reads* phenomenon can occur in this isolation level, as discussed below). As in the previous level, *range-locks* are not managed.

Putting it in simpler words, read committed is an isolation level that guarantees that any data read is committed at the moment it is read. It simply restricts the reader from seeing any intermediate, uncommitted, 'dirty' read. It makes no promise whatsoever that if the transaction re-issues the read, it will find the same data; data is free to change after it is read.

## Read uncommitted

This is the *lowest* isolation level. In this level, *dirty reads* are allowed, so one transaction may see *not-yet-committed* changes made by other transactions.

Since each isolation level is stronger than those below, in that no higher isolation level allows an action forbidden by a lower one, the standard permits a DBMS to run a transaction at an isolation level stronger than that requested (e.g., a "Read committed" transaction may actually be performed at a "Repeatable read" isolation level).

## Default isolation level

The *default isolation level* of different DBMS's varies quite widely. Most databases that feature transactions allow the user to set any isolation level. Some DBMS's also require additional syntax when performing a SELECT statement to acquire locks (e.g. *SELECT ... FOR UPDATE* to acquire exclusive write locks on accessed rows).

However, the definitions above have been criticized <sup>[3]</sup> as being ambiguous, and as not accurately reflecting the isolation provided by many databases:

This paper shows a number of weaknesses in the anomaly approach to defining isolation levels. The three ANSI phenomena are ambiguous. Even their broadest interpretations do not exclude anomalous behavior. This leads to some counter-intuitive results. In particular, lock-based isolation levels have different characteristics than their ANSI equivalents. This is disconcerting because commercial database systems typically use locking. Additionally, the ANSI phenomena do not distinguish among several isolation levels popular in commercial systems.

There are also other criticisms concerning ANSI SQL's isolation definition, in that it encourages implementors to do "bad things":

... it relies in subtle ways on an assumption that a locking schema is used for concurrency control, as opposed to an optimistic or multi-version concurrency scheme. This implies that the proposed

semantics are *ill-defined*.<sup>[4]</sup>

## Read phenomena

The ANSI/ISO standard SQL 92 refers to three different *read phenomena* when Transaction 1 reads data that Transaction 2 might have changed.

In the following examples, two transactions take place. In the first, Query 1 is performed. Then, in the second transaction, Query 2 is performed and committed. Finally, in the first transaction, Query 1 is performed again.

The queries use the following data table:

users		
id	name	age
1	Joe	20
2	Jill	25

### Dirty reads

A *dirty read* (aka *uncommitted dependency*) occurs when a transaction is allowed to read data from a row that has been modified by another running transaction and not yet committed.

Dirty reads work similarly to non-repeatable reads; however, the second transaction would not need to be committed for the first query to return a different result. The only thing that may be prevented in the READ UNCOMMITTED isolation level is updates appearing out of order in the results; that is, earlier updates will always appear in a result set before later updates.

In our example, Transaction 2 changes a row, but does not commit the changes. Transaction 1 then reads the uncommitted data. Now if Transaction 2 rolls back its changes (already read by Transaction 1) or updates different changes to the database, then the view of the data may be wrong in the records of Transaction 1.

#### Transaction 1

#### Transaction 2

```
/* Query 1 */  
SELECT age FROM users WHERE id = 1;  
/* will read 20 */
```

```
/* Query 2 */  
UPDATE users SET age = 21 WHERE id = 1;  
/* No commit here */
```

```
/* Query 1 */  
SELECT age FROM users WHERE id = 1;  
/* will read 21 */
```

```
ROLLBACK; /* lock-based DIRTY READ */
```

But in this case no row exists that has an id of 1 and an age of 21.

## Non-repeatable reads

A *non-repeatable read* occurs, when during the course of a transaction, a row is retrieved twice and the values within the row differ between reads.

*Non-repeatable reads* phenomenon may occur in a lock-based concurrency control method when read locks are not acquired when performing a SELECT, or when the acquired locks on affected rows are released as soon as the SELECT operation is performed. Under the multiversion concurrency control method, *non-repeatable reads* may occur when the requirement that a transaction affected by a commit conflict must roll back is relaxed.

### Transaction 1

```
/* Query 1 */
SELECT * FROM users WHERE id = 1;
```

### Transaction 2

```
/* Query 2 */
UPDATE users SET age = 21 WHERE id = 1;
COMMIT; /* in multiversion concurrency
         control, or lock-based READ COMMITTED */
```

```
/* Query 1 */
SELECT * FROM users WHERE id = 1;
COMMIT; /* lock-based REPEATABLE READ */
```

In this example, Transaction 2 commits successfully, which means that its changes to the row with id 1 should become visible. However, Transaction 1 has already seen a different value for *age* in that row. At the **SERIALIZABLE** and **REPEATABLE READ** isolation levels, the DBMS must return the old value for the second SELECT. At **READ COMMITTED** and **READ UNCOMMITTED**, the DBMS may return the updated value; this is a non-repeatable read.

There are two basic strategies used to prevent non-repeatable reads. The first is to delay the execution of Transaction 2 until Transaction 1 has committed or rolled back. This method is used when locking is used, and produces the serial schedule **T1, T2**. A serial schedule exhibits *repeatable reads* behaviour.

In the other strategy, as used in *multiversion concurrency control*, Transaction 2 is permitted to commit first, which provides for better concurrency. However, Transaction 1, which commenced prior to Transaction 2, must continue to operate on a past version of the database — a snapshot of the moment it

was started. When Transaction 1 eventually tries to commit, the DBMS checks if the result of committing Transaction 1 would be equivalent to the schedule **T1, T2**. If it is, then Transaction 1 can proceed. If it cannot be seen to be equivalent, however, Transaction 1 must roll back with a serialization failure.

Using a lock-based concurrency control method, at the REPEATABLE READ isolation mode, the row with ID = 1 would be locked, thus blocking Query 2 until the first transaction was committed or rolled back. In READ COMMITTED mode, the second time Query 1 was executed, the age would have changed.

Under multiversion concurrency control, at the SERIALIZABLE isolation level, both SELECT queries see a snapshot of the database taken at the start of Transaction 1. Therefore, they return the same data. However, if Transaction 1 then attempted to UPDATE that row as well, a serialization failure would occur and Transaction 1 would be forced to roll back.

At the READ COMMITTED isolation level, each query sees a snapshot of the database taken at the start of each query. Therefore, they each see different data for the updated row. No serialization failure is possible in this mode (because no promise of serializability is made), and Transaction 1 will not have to be retried.

## Phantom reads

A *phantom read* occurs when, in the course of a transaction, two identical queries are executed, and the collection of rows returned by the second query is different from the first.

This can occur when *range locks* are not acquired on performing a *SELECT ... WHERE* operation. The *phantom reads* anomaly is a special case of *Non-repeatable reads* when Transaction 1 repeats a ranged *SELECT ... WHERE* query and, between both operations, Transaction 2 creates (i.e. INSERT) new rows (in the target table) which fulfill that *WHERE* clause.

### Transaction 1

```
/* Query 1 */
SELECT * FROM users
WHERE age BETWEEN 10 AND 30;
```

### Transaction 2

```
/* Query 2 */
INSERT INTO users(id,name,age) VALUES ( 3, 'Bob', 27 );
COMMIT;
```

```
/* Query 1 */
SELECT * FROM users
WHERE age BETWEEN 10 AND 30;
COMMIT;
```

Note that Transaction 1 executed the same query twice. If the highest level of isolation were maintained, the same set of rows should be returned both times, and indeed that is what is mandated to occur in a database operating at the SQL SERIALIZABLE isolation level. However, at the lesser isolation levels, a different set



of rows may be returned the second time.

In the **SERIALIZABLE** isolation mode, Query 1 would result in all records with age in the range 10 to 30 being locked, thus Query 2 would block until the first transaction was committed. In **REPEATABLE READ** mode, the range would not be locked, allowing the record to be inserted and the second execution of Query 1 to include the new row in its results.

## Isolation Levels, Read Phenomena and Locks

### Isolation Levels vs Read Phenomena

Isolation level	Dirty reads	Non-repeatable reads	Phantoms
Read Uncommitted	may occur	may occur	may occur
Read Committed	-	may occur	may occur
Repeatable Read	-	-	may occur
Serializable	-	-	-

Anomaly Serializable is not the same as Serializable. That is, it is necessary, but not sufficient that a Serializable schedule should be free of all three phenomena types. See [1] below.

"may occur" means that the isolation level suffers that phenomenon, while "-" means that it does not suffer it.

### Isolation Levels vs Lock Duration

In lock-based concurrency control, isolation level determines the duration that locks are held.

**"C"** - Denotes that locks are held until the transaction commits.

**"S"** - Denotes that the locks are held only during the currently executing statement. Note that if locks are released after a statement, the underlying data could be changed by another transaction before the current transaction commits, thus creating a violation.

Isolation level	Write Operation	Read Operation	Range Operation (...where...)
Read Uncommitted	S	S	S
Read Committed	C	S	S
Repeatable Read	C	C	S
Serializable	C	C	C

## See also

- Atomicity
- Consistency
- Durability
- Lock (database)
- Optimistic concurrency control

- Relational Database Management System
- Snapshot isolation

## References

1. "Isolation Levels in the Database Engine", Technet, Microsoft, [http://technet.microsoft.com/en-us/library/ms189122\(v=SQL.105\).aspx](http://technet.microsoft.com/en-us/library/ms189122(v=SQL.105).aspx)
2. "The Architecture of Transaction Processing Systems", Chapter 23, Evolution of Processing Systems, Department of Computer Science, Stony Brook University, retrieved 20 March 2014, <http://www.cs.sunysb.edu/~liu/cse315/23.pdf>
3. "A Critique of ANSI SQL Isolation Levels" (<http://www.cs.umb.edu/~poneil/iso.pdf>) (PDF). Retrieved 29 July 2012.
4. salesforce (2010-12-06). "Customer testimonials (SimpleGeo, CLOUDSTOCK 2010)" (<http://www.youtube.com/v/7J61pPG9j90?version=3>). www.DataStax.com: DataStax. Retrieved 2010-03-09. " (see above at about 13:30 minutes of the webcast!)"

## External links

- Oracle® Database Concepts ([http://docs.oracle.com/cd/B12037\\_01/server.101/b10743/toc.htm](http://docs.oracle.com/cd/B12037_01/server.101/b10743/toc.htm)), chapter 13 Data Concurrency and Consistency, Preventable Phenomena and Transaction Isolation Levels ([http://docs.oracle.com/cd/B12037\\_01/server.101/b10743/consist.htm#sthref1919](http://docs.oracle.com/cd/B12037_01/server.101/b10743/consist.htm#sthref1919))
- Oracle® Database SQL Reference ([http://docs.oracle.com/cd/B19306\\_01/server.102/b14200/toc.htm](http://docs.oracle.com/cd/B19306_01/server.102/b14200/toc.htm)), chapter 19 SQL Statements: SAVEPOINT to UPDATE ([http://docs.oracle.com/cd/B19306\\_01/server.102/b14200/statements\\_10.htm#i2068385](http://docs.oracle.com/cd/B19306_01/server.102/b14200/statements_10.htm#i2068385)), SET TRANSACTION ([http://docs.oracle.com/cd/B19306\\_01/server.102/b14200/statements\\_10005.htm#i2067247](http://docs.oracle.com/cd/B19306_01/server.102/b14200/statements_10005.htm#i2067247))
- in JDBC: Connection constant fields ([http://docs.oracle.com/javase/7/docs/api/java/sql/Connection.html#field\\_summary](http://docs.oracle.com/javase/7/docs/api/java/sql/Connection.html#field_summary)), `Connection.getTransactionIsolation()` ([http://docs.oracle.com/javase/7/docs/api/java/sql/Connection.html#getTransactionIsolation\(\)](http://docs.oracle.com/javase/7/docs/api/java/sql/Connection.html#getTransactionIsolation())), `Connection.setTransactionIsolation(int)` ([http://docs.oracle.com/javase/7/docs/api/java/sql/Connection.html#setTransactionIsolation\(int\)](http://docs.oracle.com/javase/7/docs/api/java/sql/Connection.html#setTransactionIsolation(int)))
- in Spring Framework: `@Transactional` (<http://static.springsource.org/spring/docs/current/javadoc-api/org/springframework/transaction/annotation/Transactional.html>), Isolation (<http://static.springsource.org/spring/docs/current/javadoc-api/org/springframework/transaction/annotation/Isolation.html>)

Retrieved from "https://en.wikipedia.org/w/index.php?title=Isolation\_(database\_systems)&oldid=673395169"

Categories: Data management | Transaction processing

- 
- This page was last modified on 28 July 2015, at 01:13.
  - Text is available under the Creative Commons Attribution-ShareAlike License; additional terms may apply. By using this site, you agree to the Terms of Use and Privacy Policy. Wikipedia® is a registered trademark of the Wikimedia Foundation, Inc., a non-profit organization.

# Java Information

Monday, May 23, 2011

## Dirty Read,Phantom Read and Non Repeatable Read

### What is Dirty Read?

Dirty read occurs wherein one transaction is changing the tuple/record, and a second transaction can read this tuple/record before the original change has been committed or rolled back. This is known as a dirty read scenario because there is always the possibility that the first transaction may rollback the change, resulting in the second transaction having read an invalid value.

To understand it better,lets take a use case where on thread is viewing the record and other thread is updating the value of the record.Since the transaction isolation attribute is set to READ UNCOMMITTED.Second thread will be able to see the changes made by other thread.

#### Example of Dirty Read:-

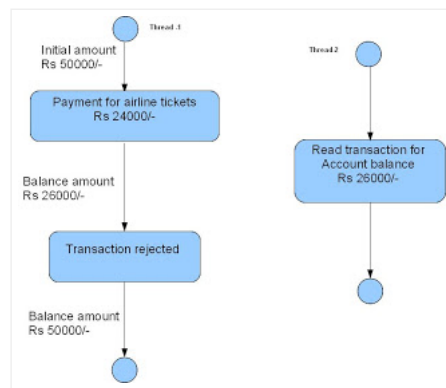
In the current use case we have a table containing account balances. One thread is reading the data from the account balances table.Other thread is updating the data from that table.Since isolation attribute is 'READ UNCOMMITTED' .Other thread can view non committed data.

Table Definition:-

Create table AccountBalance

```
(
    id integer Primary Key,
    acctName varchar2(255) not null,
    acctBalance integer(9,2) not null,
    bankName varcahr2(255)
);
```

```
insert into AccountBalance values (1,'Kunaal',50000,'Bank-a');
```



ReaderRunImpl.java

```
package com.kunaal.pooling;

import java.sql.Connection;
import java.sql.PreparedStatement;
import java.sql.ResultSet;
import java.sql.SQLException;

/**
 * Reader thread which selects the data while other thread is updating the
 * data
 *
 * @author Kunaal A Trehan
 *
 */
public class ReaderRunImpl implements Runnable{

    private Connection conn;

    private static final String QUERY="Select balance from AccountBalance where id=1";
```

#### Labels

- [auto boxing](#)
- [auto unboxi](#)
- [Enums \(1\)](#)
- [JUnit4 \(1\)](#)
- [Mockito \(1\)](#)
- [Overloading](#)
- [paramterize Java \(1\)](#)
- [plain enum](#)
- [PowerMock](#)

#### Blog Archive


- ▼ [2011 \(18\)](#)
  - [Feb 2011](#)
  - [Mar 2011](#)
  - ▼ [May 2011](#)
    - Dirty Rea
    - Read .
    - Repee
  - [Jul 2011](#)
  - [Aug 2011](#)
  - [Sep 2011](#)
  - [Nov 2011](#)
  - [Dec 2011](#)

- [2012 \(2\)](#)
- [2013 \(1\)](#)

#### Followers

Join th  
with Google Fr

#### Members (1)



Already a mem

#### About Me

**Kunaal A 1**



[View my comp](#)

#### Subscribe To



```

public ReaderRunImpl(Connection conn){
    this.conn=conn;
}

@Override
public void run() {
    PreparedStatement stmt =null;
    ResultSet rs =null;

    try {
        stmt = conn.prepareStatement(QUERY);
        rs = stmt.executeQuery();
        while (rs.next()){
            System.out.println("Balance is:" + rs.getDouble(1));
        }
    } catch (SQLException e) {
        e.printStackTrace();
    }finally{
        try {
            stmt.close();
            rs.close();
        } catch (SQLException e) {
            e.printStackTrace();
        }
    }
}
}

```

#### PaymentRunImpl.java

```

package com.kunaal.pooling;

import java.sql.Connection;
import java.sql.PreparedStatement;
import java.sql.SQLException;

/**
 * @author Kunaal A Trehan
 */
public class PaymentRunImpl implements Runnable{

    private Connection conn;

    private static final String QUERY="Update AccountBalance set balance=26000 where id=1";

    public PaymentRunImpl(Connection conn){
        this.conn=conn;
    }

    @Override
    public void run() {
        PreparedStatement stmt = null;

        try {
            stmt = conn.prepareStatement(QUERY);
            stmt.execute();
            Thread.currentThread().sleep(3000);
            conn.rollback();
        } catch (SQLException e) {
            e.printStackTrace();
        } catch (InterruptedException e) {
            e.printStackTrace();
        }finally{
            try {
                stmt.close();
            } catch (SQLException e) {
                e.printStackTrace();
            }
        }
    }
}

```

#### DirtyReadExample.java

```

package com.kunaal.pooling;

import java.sql.Connection;
import java.sql.SQLException;

/**
 * @author Kunaal A Trehan
 */
public class DirtyReadExample {

```

```

/**
 * @param args
 * @throws SQLException
 * @throws InterruptedException
 */
public static void main(String[] args) {
    ConnectionPool pool=new ConnectionPool(5,1000);

    Connection connPymt = pool.getConnection();
    Connection connReader = pool.getConnection();
    try {
        connPymt.setAutoCommit(false);
        connPymt.setTransactionIsolation(Connection.TRANSACTION_READ_COMMITTED);

        connReader.setAutoCommit(false);
        //connReader.setTransactionIsolation(Connection.TRANSACTION_READ_UNCOMMITTED);
        connReader.setTransactionIsolation(Connection.TRANSACTION_READ_COMMITTED);
    } catch (SQLException e) {
        e.printStackTrace();
    }

    Thread pymtThread=new Thread(new PaymentRunImpl(connPymt));
    Thread readerThread=new Thread(new ReaderRunImpl(connReader));

    pymtThread.start();
    try {
        Thread.currentThread().sleep(1000);
    } catch (InterruptedException e) {
        // TODO Auto-generated catch block
        e.printStackTrace();
    }
    readerThread.start();

    pool.returnConnection(connPymt);
    pool.returnConnection(connReader);
}
}

```

Here reader thread views the account balance when payment thread is sleeping. So reader thread will view the balance as 24000 if the isolation level is Connection.TRANSACTION\_READ\_UNCOMMITTED as other transaction has roll back the transaction.

## What is Phantom Read?

Phantom read occurs where in a transaction same query executes twice, and the second result set includes rows that weren't visible in the first result set. This situation is caused by another transaction inserting new rows between the execution of the two queries.

### Example of Phantom Read:-

To understand it better consider a use case where one thread is inserting the data while other thread is reading the data in different transaction. Since reader thread has isolation attribute as READ COMMITTED, reader thread will see the new rows inserted when it queries again.

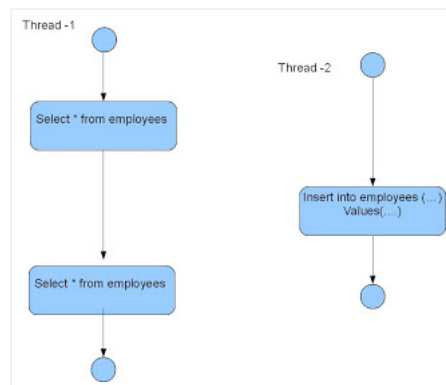
Table definition:-

Create table Employee

```

(
    id integer Primary Key,
    empName varchar2(255) not null,
    empCity varchar2(255) not null,
    empCtry varchar2(255) not null
);

```



PhantomReader.java

```

package com.kunaal.pooling;

import java.sql.Connection;

```

```

import java.sql.PreparedStatement;
import java.sql.ResultSet;
import java.sql.SQLException;

/**
 * @author Kunaal A Trehan
 *
 */
public class PhantomReader implements Runnable{

    private Connection conn;

    private static final String QUERY="Select * from Employee";

    public PhantomReader(Connection conn){
        this.conn=conn;
    }

    @Override
    public void run() {
        PreparedStatement stmt =null;
        ResultSet rs=null;

        try {
            stmt=conn.prepareStatement(QUERY);
            rs=stmt.executeQuery();

            while(rs.next()){
                System.out.println("Emp Details- " + rs.getInt(1) + "-" + rs.getString(2) + "-" +
                    rs.getString(3) + "-" + rs.getString(4));
            }

            Thread.currentThread().sleep(3000);
            System.out.println("AFTER WAKING UP");
            System.out.println("=====");

            rs=stmt.executeQuery();

            while(rs.next()){
                System.out.println("Emp Details- " + rs.getInt(1) + "-" + rs.getString(2) + "-" +
                    rs.getString(3) + "-" + rs.getString(4));
            }
        } catch (SQLException e) {
            e.printStackTrace();
        } catch (InterruptedException e) {
            e.printStackTrace();
        } finally{
            try {
                rs.close();
                stmt.close();
            } catch (SQLException e) {
                e.printStackTrace();
            }
        }
    }
}

```

PhantomInsert.java

```

package com.kunaal.pooling;

import java.sql.Connection;
import java.sql.PreparedStatement;
import java.sql.SQLException;

/**
 * @author Kunaal A Trehan
 *
 */
public class PhantomInsert implements Runnable{

    private Connection conn;

    private static final String QUERY="Insert into Employee values(?,?,?,?)";

    public PhantomInsert(Connection conn){
        this.conn=conn;
    }

    @Override
    public void run() {
        PreparedStatement stmt =null;

        try {
            stmt = conn.prepareStatement(QUERY);
            stmt.setInt(1, 3);

```

```

stmt.setString(2, "ABC");
stmt.setString(3, "DELHI");
stmt.setString(4, "INDIA");

stmt.execute();
conn.commit();
} catch (SQLException e) {
    e.printStackTrace();
} finally {
    try {
        stmt.close();
    } catch (SQLException e) {
        e.printStackTrace();
    }
}
}
}

```

PhantomReadExample.java

```

package com.kunaal.pooling;

import java.sql.Connection;
import java.sql.SQLException;

/**
 * @author Kunaal A Trehan
 */
public class PhantomReadExample {

    /**
     * @param args
     */
    public static void main(String[] args) {
        ConnectionPool pool=new ConnectionPool(5,1000);

        Connection connInsert = pool.getConnection();
        Connection connReader = pool.getConnection();
        try {
            connInsert.setAutoCommit(false);
            connInsert.setTransactionIsolation(Connection.TRANSACTION_READ_COMMITTED);

            connReader.setAutoCommit(false);
            //connInsert.setTransactionIsolation(Connection.TRANSACTION_READ_COMMITTED);
            connReader.setTransactionIsolation(Connection.TRANSACTION_SERIALIZABLE);
        } catch (SQLException e) {
            e.printStackTrace();
        }

        Thread readThread=new Thread(new PhantomReader(connReader));
        Thread insertThread=new Thread(new PhantomInsert(connInsert));
        readThread.start();
        insertThread.start();

        pool.returnConnection(connReader);
        pool.returnConnection(connInsert);

    }
}

```

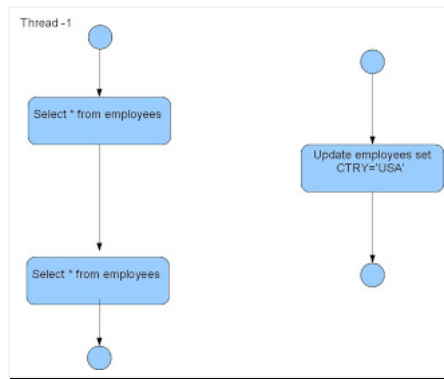
Here reader thread with isolation level as READ COMMITTED queries the employee table twice while other thread is inserting the data. As a result of which number of rows returned is different.

## What is Non Repeatable Read?

Non Repeatable Reads happen when in a same transaction same query yields different results. This happens when another transaction updates the data returned by other transaction.

### Example of Non Repeatable Read:-

To understand it better let's take a use case where one thread is viewing the data and other thread is updating the data. Since isolation level is READ COMMITTED, other thread will be able to view the updated changes. So in the same transaction, same query will yield different data.



Updater.java

```

package com.kunaal.pooling;

import java.sql.Connection;
import java.sql.PreparedStatement;
import java.sql.SQLException;

/**
 * @author Kunaal A Trehan
 */
public class Updater implements Runnable{

    private Connection conn;

    private static final String QUERY="Update Employee set empCountry='USA'";

    public Updater(Connection conn){
        this.conn=conn;
    }

    @Override
    public void run() {
        PreparedStatement stmt =null;

        try {
            stmt = conn.prepareStatement(QUERY);
            stmt.executeUpdate();
            conn.commit();
        } catch (SQLException e) {
            e.printStackTrace();
        } finally{
            try {
                stmt.close();
            } catch (SQLException e) {
                e.printStackTrace();
            }
        }
    }
}

```

NonRepeatableReader.java

```

package com.kunaal.pooling;

import java.sql.Connection;
import java.sql.PreparedStatement;
import java.sql.ResultSet;
import java.sql.SQLException;

/**
 * @author Kunaal A Trehan
 */
public class NonRepeatableReader implements Runnable{

    private Connection conn;

    private static final String QUERY="Select * from Employee";

    public NonRepeatableReader (Connection conn){
        this.conn=conn;
    }

    @Override

```



```

public void run() {
    PreparedStatement stmt =null;
    ResultSet rs=null;

    try {
        stmt=conn.prepareStatement (QUERY);
        rs=stmt.executeQuery();

        while(rs.next()){
            System.out.println("Emp Details- " + rs.getInt(1) + "-" + rs.getString(2) + "-" +
                rs.getString(3) + "-" + rs.getString(4));
        }

        Thread.currentThread().sleep(3000);
        System.out.println("AFTER WAKING UP");
        System.out.println("=====");

        rs=stmt.executeQuery();

        while(rs.next()){
            System.out.println("Emp Details- " + rs.getInt(1) + "-" + rs.getString(2) + "-" +
                rs.getString(3) + "-" + rs.getString(4));
        }
    } catch (SQLException e) {
        e.printStackTrace();
    } catch (InterruptedException e) {
        e.printStackTrace();
    } finally{
        try {
            rs.close();
            stmt.close();
        } catch (SQLException e) {
            e.printStackTrace();
        }
    }
}
}

```

NonRepeatbleExample.java

```

package com.kunaal.pooling;

import java.sql.Connection;
import java.sql.SQLException;

/**
 * @author Kunaal A Trehan
 */
public class NonRepeatbleExample {

    /**
     * @param args
     */
    public static void main(String[] args) {
        ConnectionPool pool=new ConnectionPool(5,1000);

        Connection connUpdt = pool.getConnection();
        Connection connReader = pool.getConnection();
        try {
            connUpdt.setAutoCommit(false);
            connUpdt.setTransactionIsolation(Connection.TRANSACTION_READ_COMMITTED);

            connReader.setAutoCommit(false);
            //connReader.setTransactionIsolation(Connection.TRANSACTION_READ_UNCOMMITTED);
            connReader.setTransactionIsolation(Connection.TRANSACTION_READ_COMMITTED);
        } catch (SQLException e) {
            e.printStackTrace();
        }

        Thread updtThread=new Thread(new Updater(connUpdt));
        Thread readerThread=new Thread(new NonRepeatableReader(connReader));

        readerThread.start();
        try {
            Thread.currentThread().sleep(1000);
        } catch (InterruptedException e) {
            // TODO Auto-generated catch block
            e.printStackTrace();
        }
        updtThread.start();

        pool.returnConnection(connUpdt);
        pool.returnConnection(connReader);
    }
}

```

Here while reading same query will yield different results as other thread has updated the data.

Following table depicts the isolation level mapping with dirty read ,phantom read and others.

ISOLATION LEVEL	Dirty Read	Non Repeatable Read	Phantom Read
READ_COMMITTED	YES	NO	NO
REPEATABLE_READ	YES	YES	NO
SERIALIZABLE	YES	YES	YES

Posted by Kunaal A Trehan at 7:24 AM



+2 Recommend this on Google

## 9 comments:



**Anonymous** June 15, 2012 at 7:55 AM

Where is this ConnectionPool class came from?  
Is this something that you designed?

I could not able to find any libraries for this class.

[Reply](#)

[Replies](#)



**Sergey** December 5, 2012 at 12:39 PM

it is really does not matter, actually

[Reply](#)



**CaR** August 4, 2012 at 9:03 AM

thanks, very good information, I wasn't sure about NR read VS phantom read, now I am.

[Reply](#)



**Nitesh Porwal** November 15, 2013 at 7:47 PM

thanks, cleared the conception of Isolation level in transaction.

[Reply](#)



**Алексей Кузнецов** June 22, 2014 at 3:48 AM

In example of Phantom read you can change isolation level to REPEATABLE\_READ, so Phantom read problem remain.  
But when you start programm you will see that there is no Phantom record! Im confused why does it happen

[Reply](#)



**Anonymous** [November 8, 2014 at 6:14 AM](#)

Good job !

[Reply](#)



**Manish** [March 25, 2015 at 12:56 AM](#)

Nice post

[Reply](#)



**Anonymous** [May 10, 2015 at 11:53 PM](#)

good job

[Reply](#)

[Replies](#)



**Anonymous** [July 25, 2015 at 9:45 PM](#)

nice explanation

---

[Reply](#)

Enter your comment...

Comment as: Google Account ▼

[Publish](#)

[Preview](#)

## Links to this post

[Create a Link](#)

[Newer Post](#)

[Home](#)

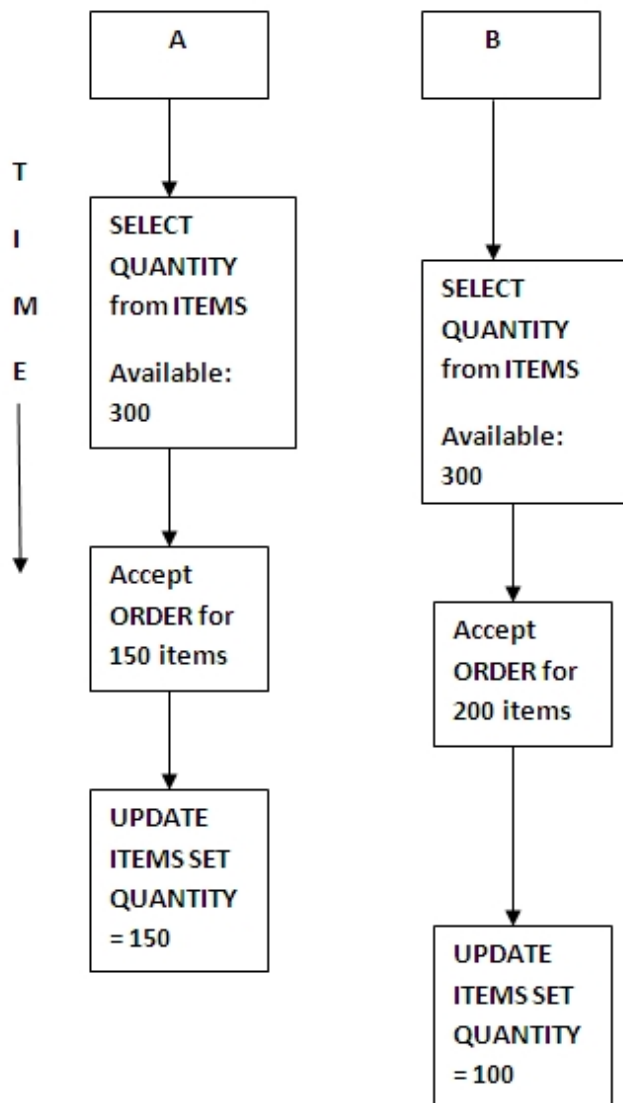
[Older Post](#)

Subscribe to: [Post Comments \(Atom\)](#)

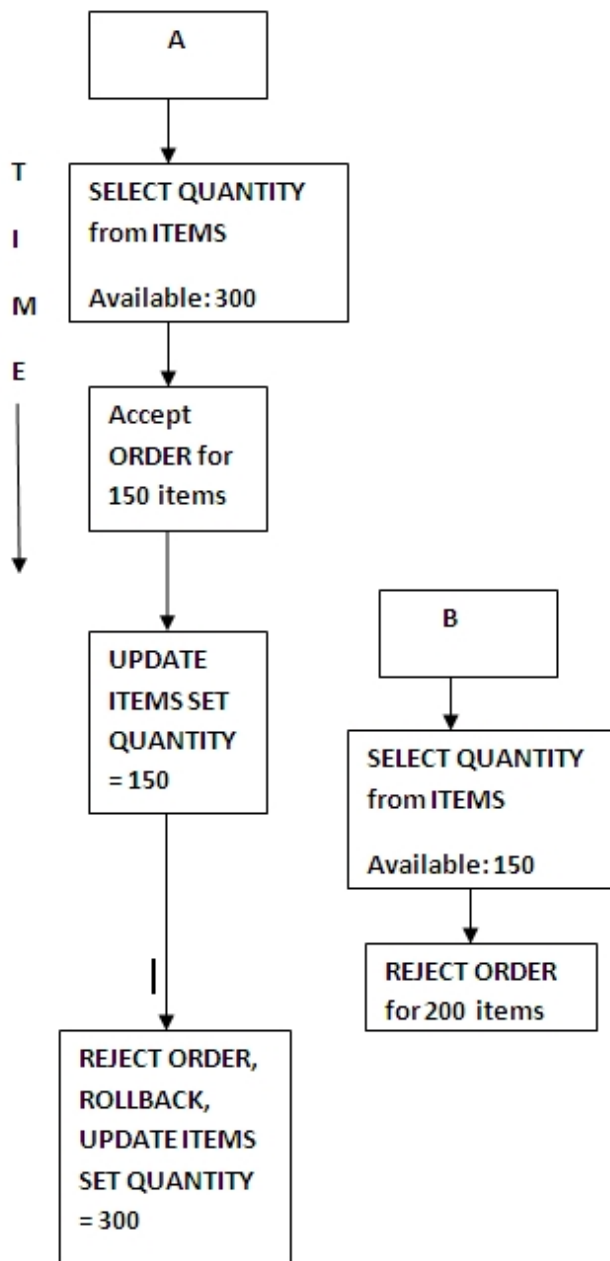
# Lost update, uncommitted data, dirty read problem in transaction processing in SQL

March 14, 2013 by Admin — [Leave a Comment](#)

**Lost update problem:** A lost update is a typical problem in transaction processing in SQL. It happens when two queries access and update the same data from a database. This problem can be understood by the below given diagram. Here, A is processing an order for a client for 150 items. He checks from the ITEMS table that there are 300 available items. So he starts placing the order. After few seconds B gets an order for 200 items. He also checks the items table and finds that there are 300 items available. So he also starts placing an order. Meanwhile A confirms the order for 150 items and updates the ITEMS table and sets the quantity to 150. A few seconds later B confirms the order and updates the ITEMS table and sets the quantity to 100. This problem is known as lost update problem because both the orders of the users A and B have been accepted, but there is not enough items available. Hence, the updates are lost.



**Uncommitted data or dirty read problem:** In a dirty read problem, A is processing an order for a client for 150 items. He checks from the ITEMS table that there are 300 available items. So he starts placing the order. A confirms the order for 150 items and updates the ITEMS table and sets the quantity to 150. Now, B receives an order for 200 items. He checks the ITEMS table to find that there is not enough inventories (150 available) and rejects the order. By using business rules and transactions a note is sent informing that more items are required. Now, due to some reason client asks A to cancel the order, so A cancels the order, rolls back and updates the ITEMS table back to 300 items. This problem is known as dirty read because B saw the uncommitted update of A.



#### Share this:



#### Related

Inconsistent data, non repeatable read, phantom insert problem in transaction processing in SQL

March 15, 2013

In "SQL"

Versioning in multiuser transaction processing in SQL: advantages, disadvantages

March 18, 2013

In "SQL"

Isolation levels in SQL: serializable, repeatable read, committed, uncommitted, SET TRANSACTION statement

March 17, 2013

In "SQL"

Filed Under: [SQL](#)

Tagged With: [data](#), [dirty](#), [in](#), [Lost](#), [problem](#), [processing](#), [read](#), [SQL](#), [transaction](#), [uncommitted](#), [update](#)

---

## Chromebooks for Work



Built Beautiful, Built for Work Live Online Event - Register Now!



---

## Leave a Reply

Enter your comment here...

Search the site ...

## Application Performance Management



APP SERVERS,  
SERVERS, DATABASES,  
CLOUD SERVICES, SAP,  
VIRTUAL SERVERS, etc ..

**ManageEngine**  
Applications Manager



**DOWNLOAD  
FREE TRIAL**

RECENT POSTS

CATEGORIES

Virtualization: The Advantages and Disadvantages of VM  
April 27, 2015

Internal Security Control Classifications April 24, 2015

Apache Archiva Cross-Site scripting and Command  
Execution April 22, 2015

MongoDB Accessible Without Authentication Vulnerability  
April 20, 2015

Google Chrome Prior to 42.0.2311.90 Multiple  
Vulnerabilities April 16, 2015

## Categories

Select Category ▼

## ARCHIVES

### Archives

Select Month ▼

## FOLLOW US



## SUBSCRIBE VIA RSS





## SUBSCRIBE VIA EMAIL

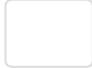
Enter your email address to subscribe to receive  
notifications of new posts by email.

SUBSCRIBE



 **This summer, meet online from anywhere.** 

**Try it free**

MENU 

## Spring transaction isolation level tutorial

30 January 2013

By [Gonçalo Marques](#)

java spring tx

In this tutorial you will learn about the transaction isolation level provided by the Spring framework.

### Introduction

Transaction isolation level is a concept that is not exclusive to the Spring framework. It is applied to transactions in general and is directly related with the ACID transaction properties. Isolation level defines how the changes made to some data repository by one transaction affect other simultaneous concurrent transactions, and also how and when that changed data becomes available to other transactions. When we define a transaction using the Spring framework we are also able to configure in which isolation level that same transaction will be executed.

### Usage example

Using the **@Transactional** annotation we can define the isolation level of a Spring managed bean transactional method. This means that the transaction in which this method is executed will run with that isolation level:

#### Isolation level in a transactional method

```
@Autowired
private TestDAO testDAO;

@Transactional(isolation=Isolation.READ_COMMITTED)
public void someTransactionalMethod(User user) {

    // Interact with testDAO


}
```

We are defining this method to be executed in a transaction which isolation level is **READ\_COMMITTED**. We will see each isolation level in detail in the next sections.

# READ\_UNCOMMITTED

**READ\_UNCOMMITTED** isolation level states that a transaction **may** read data that is still **uncommitted** by other transactions. This constraint is very relaxed in what matters to transactional concurrency but it may lead to some issues like **dirty reads**. Let's see the following image:

Dirty read



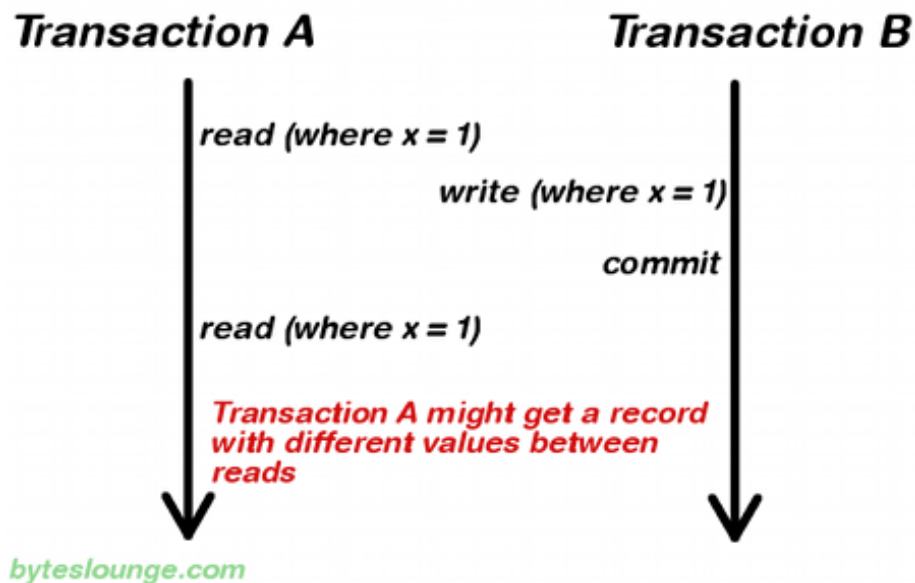
In this example **Transaction A** writes a record. Meanwhile **Transaction B** reads that same record before **Transaction A** commits. Later **Transaction A** decides to rollback and now we have changes in **Transaction B** that are inconsistent. This is a **dirty read**. **Transaction B** was running in **READ\_UNCOMMITTED** isolation level so it was able to read **Transaction A** changes before a commit occurred.

**Note:** **READ\_UNCOMMITTED** is also vulnerable to **non-repeatable reads** and **phantom reads**. We will also see these cases in detail in the next sections.

# READ\_COMMITTED

**READ\_COMMITTED** isolation level states that a transaction can't read data that is **not** yet committed by other transactions. This means that the **dirty read** is no longer an issue, but even this way other issues may occur. Let's see the following image:

Non-repeatable read



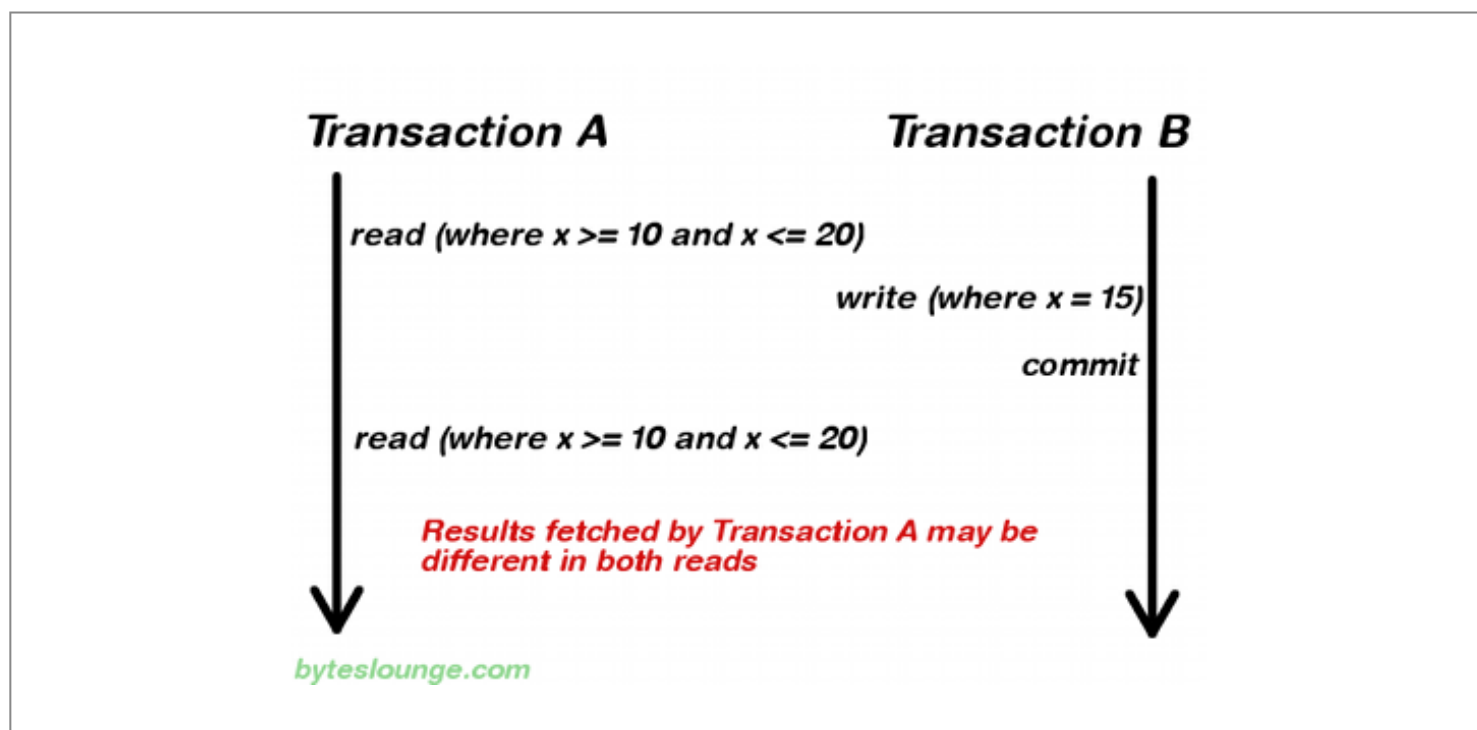
In this example **Transaction A** reads some record. Then **Transaction B** writes that same record and commits. Later **Transaction A** reads that same record again and may get different values because **Transaction B** made changes to that record and committed. This is a **non-repeatable read**.

**Note:** **READ\_COMMITTED** is also vulnerable to **phantom reads**. We will also see this case in detail in the next section.

# REPEATABLE\_READ

**REPEATABLE\_READ** isolation level states that if a transaction reads one record from the database multiple times the result of all those reading operations must always be the same. This eliminates both the **dirty read** and the **non-repeatable read** issues, but even this way other issues may occur. Let's see the following image:

Phantom read



In this example **Transaction A** reads a **range** of records. Meanwhile **Transaction B** inserts a new record in the same range that **Transaction A** initially fetched and commits. Later **Transaction A** reads the same range again and will also get the record that **Transaction B** just inserted. This is a **phantom read**: a transaction fetched a range of records multiple times from the database and obtained different result sets (containing phantom records).

## SERIALIZABLE

**SERIALIZABLE** isolation level is the most restrictive of all isolation levels. Transactions executed with this level

**SERIALIZABLE** isolation level is the most restrictive of all isolation levels. Transactions are executed with locking at all levels (**read**, **range** and **write** locking) so they appear as if they were executed in a serialized way. This leads to a scenario where **none** of the issues mentioned above may occur, but in the other way we don't allow transaction concurrency and consequently introduce a performance penalty.

## DEFAULT

**DEFAULT** isolation level, as the name states, uses the default isolation level of the datastore we are actually connecting from our application.

## Summary

To summarize, the existing relationship between isolation level and read phenomena may be expressed in the following table:

	dirty reads	non-repeatable reads	phantom reads
READ_UNCOMMITTED	yes	yes	yes
READ_COMMITTED	no	yes	yes
REPEATABLE_READ	no	no	yes
SERIALIZABLE	no	no	no

## JPA

If you are using Spring with JPA you may come across the following exception when you use an isolation level that is different the default:

```
InvalidIsolationLevelException: Standard JPA does not support custom isolation levels - use a special
JpaDialect for your JPA implementation
at org.springframework.orm.jpa.DefaultJpaDialect.beginTransaction(DefaultJpaDialect.java:67)
at org.springframework.orm.jpa.JpaTransactionManager.doBegin(JpaTransactionManager.java:378)
at
org.springframework.transaction.support.AbstractPlatformTransactionManager.getTransaction(AbstractPlatform
at
org.springframework.transaction.interceptor.TransactionAspectSupport.createTransactionIfNecessary(Transaction
at
org.springframework.transaction.interceptor.TransactionAspectSupport.invokeWithinTransaction(TransactionAs
at org.springframework.transaction.interceptor.TransactionInterceptor.invoke(TransactionInterceptor.java:94)
at
org.springframework.aop.framework.ReflectiveMethodInvocation.proceed(ReflectiveMethodInvocation.java:172)
at org.springframework.aop.framework.JdkDynamicAopProxy.invoke(JdkDynamicAopProxy.java:204)
```

To solve this problem you must implement a custom JPA dialect which is explained in detail in the following article: [Spring - Change transaction isolation level example](#).



## Related Articles

- [Spring JDBC transactions example](#)
- [Spring with Hibernate persistence and transactions example](#)
- [Spring transaction propagation tutorial](#)
- [Spring JTA multiple resource transactions in Tomcat with Atomikos example](#)

## Comments

Post Comment

**Francisco A. Lozano**

31 January 2013

It's important to mention that setting a the right isolation level in the datasource or datastore's defaults instead of forcing to explicit isolation level in annotations, together with the right connection string parameters in MySQL (so that it doesn't change always isolation level) can boost performance quite a lot

Reply

**gmarques**

31 January 2013

Hello, Thank you for your feedback. Can you provide reference documentation on that assumption?

Reply

**Kishore**

22 April 2014

Dear Gonalo Marques I help me a lot to understand how spring supports Isolation level. I hope i helps so many. Thank You very much.

Reply

**Mohd Kose Avase**

31 July 2014

Good explanation. Thank you very much

Reply

**TechyKudos**

16 March 2015

It is really helpful to me. Very nice to understand.

Thanks a lot for your efforts, really appreciable.

Reply

**Felipe Gutierrez**

17 March 2015

Very good explanation. Thanks for that!

Reply

## About the author

Gonçalo Marques is a Software Engineer with several years of experience in software development and architecture definition. During this period his main focus was delivering software solutions in banking, telecommunications and governmental areas. He created the Bytes Lounge website with one ultimate goal: share his knowledge with the software development community. His main area of expertise is Java and open source.

GitHub profile: <http://github.com/gonmarques>

He is also the author of the **WiFi File Browser** Android application:

## SOCIAL

354

Like

78

Follow

24

g+1

## ABOUT



Gonçalo Marques




Follow

140

Java and Open Source enthusiast

[GitHub profile](#)

DON'T BE SURPRISED WHEN  
YOUR SERVERS FAIL.



Predict, manage & maximize  
with **Solarwinds Server  
& Application Monitor.**

solarwinds

LEARN MORE »




Bytes Lounge  
354 likes

Like Page

Share

Be the first of your friends to like this





## SUBSCRIBE

Receive the latest articles in your mailbox

Subscribe

## POPULAR TAGS

spring jpa mvc  
java-ee java jsf  
web java8 cdi

## ARCHIVE

2015 (7)

2014 (32)

2013 (55)

2012 (9)

Java Collections

Java IO

Design Patterns

Java Concurrency

Java Reflection

Java 8

CDI

EJB

JPA

Web

JSF

JAAS

Spring Core

Spring MVC

[Android](#)  
[Maven](#)  
[Postfix](#)  
[New Relic](#)